

Cross-Dataset Intrusion Detection: Problem Impact, Mathematical Diagnosis & Engineering Solutions

A Complete Guide to Resolving the Generalization Gap in Machine Learning Network Defense

Authors: Pynskhemlang Marbaniang, Tehjib Almas Junaid, Bamelaa Marbaniang, Md Sabique Huda	Department: Computer Science & Engineering
Guide: Mr. Abdul Kayum Ali Supervisor: Mr. Debarbat Bharali	Institution: University of Science & Technology, Meghalaya (USTM)

Part 1: The Core Problem & Real-World Cybersecurity Impact

Network Intrusion Detection Systems (NIDS) inspect network packets and flows to detect cyber attacks in real time. Over the last decade, academic researchers transitioned from signature-based tools (Snort, Suricata) to supervised Machine Learning (Random Forests, SVM, XGBoost). In published research, these ML classifiers routinely achieve **98% to 99.5% accuracy** on benchmark datasets like CICIDS2017 and UNSW-NB15. However, when these identical models are transferred into a live enterprise network or tested against an independent dataset, **their accuracy catastrophically collapses down to 40%–60%**.

■ HOW THIS IMPACTS REAL-WORLD CYBERSECURITY DEFENSE:

- 1. Deadly False Sense of Security:** C-level executives and security teams believe their perimeter has 99% AI protection, while in reality the system fails to detect over half of incoming attacks.
- 2. SOC Analyst Alert Fatigue:** In cross-network transfer, False Alarm Rates (FAR) surge by 300% to 500%. Security analysts receive thousands of false positive alerts daily, causing them to ignore genuine intrusions.
- 3. Zero-Day & Variant Vulnerability:** Attackers subtly alter packet payload lengths or transmission timings to bypass the model, because the model never learned true attack logic—it only memorized laboratory noise.

Part 2: Mathematical & Statistical Diagnosis (Why It Collapses)

The performance collapse is driven by three distinct statistical phenomena:

1. Covariate Shift (Feature Distribution Shift — $P(X)$ Changes):

The underlying relationship between network behavior and attack intent remains constant, but the input feature distributions differ radically. For example, UNSW-NB15 was captured on a **10 Gbps enterprise backbone** with high packet rates, while CICIDS2017 was captured on a **1 Gbps testbed**. A normal flow in UNSW may transmit 50,000 bytes per second, which an ML model trained on CICIDS misclassifies as a high-volume DoS flood.

2. Concept Drift (Label Distribution Shift — $P(Y|X)$ Changes):

Attack tools and network protocols evolve. A benign file download over HTTP in 2017 exhibits different packet burst sizes and inter-arrival times compared to synthetic HTTP traffic generated by automated fuzzers in 2015.

3. Laboratory Shortcut Learning (Artifact Overfitting):

Standard research trains and tests on 80/20 splits of a single dataset. In this closed world, the ML model cheats: it memorizes specific source IP subnets, operating system TCP window sizes, or constant port numbers instead of learning behavioral anomaly semantics.

Table 1: Structural Divergence Between the Two Benchmark Datasets

Dimension	CICIDS2017 Benchmark	UNSW-NB15 Benchmark
Generating Institution	Canadian Institute for Cybersecurity (UNB)	Australian Centre for Cyber Security (UNSW)
Extraction Toolchain	CICFlowMeter (80+ statistical flow features)	Argus & Bro-Zeek (49 flow/content features)
Physical Network Speed	1 Gbps network topology	10 Gbps high-speed IXIA PerfectStorm testbed
Attack Taxonomies	7 families (DoS, DDoS, PortScan, Brute Force, Web, Botnet, Infiltration)	9 categories (Fuzzers, Analysis, Backdoors, DoS, Exploits, Generic, Recon, Shellcode, Worms)
Class Imbalance	>80% Benign, ~20% Attack	~68% Benign, ~32% Attack

Part 3: How We Solve the Issue — Our 4-Layer Engineering Solution

Prior papers (e.g., Verkerken et al., 2022 and Cantone et al., 2024) observed that accuracy plummets by 28%–42% and stopped, calling it an open problem. **Our project engineers an active 4-layer remediation system:**

Layer 1: Semantic Feature Ontology Mapping (Harmonizing Naming & Units)

Because CICIDS uses CICFlowMeter and UNSW uses Argus/Zeek, column names and measurements differ. We build an explicit semantic translation dictionary connecting equivalent behavioral attributes into a shared feature space: Flow Duration ↔ dur, Total Fwd/Bwd Packets ↔ spkts/dpkts, Total Length of Fwd/Bwd Packets ↔ sbytes/dbytes, Flow Packets/s ↔ rate, and Packet Length Mean/Std ↔ smean/dmean. All unmapped, dataset-specific columns are discarded.

Layer 2: Statistical Feature Selection (Discarding Shortcut Features)

We apply two rigorous feature selection techniques to prune features that fail to generalize across networks:

- **Mutual Information (MI) Scoring:** Calculates non-linear information dependency between features and the attack label across both datasets. We retain only features with high predictive power in *both* environments.
- **Recursive Feature Elimination (RFE):** Iteratively trains tree ensembles and eliminates variables with high dataset-specific variance.

Layer 3: Domain-Invariant Normalization (Neutralizing 1G vs 10G Speeds)

Standard Z-score scaling (StandardScaler) fails because mean and variance are distorted by network bandwidth. We implement **Quantile Transformation** (mapping features to uniform distributions) and **Robust Z-Score** (using median and IQR). This equalizes magnitude orders across 1 Gbps and 10 Gbps networks without corrupting the underlying shape of attack anomalies.

Layer 4: CORAL (Correlation Alignment) Domain Adaptation

CORAL is an unsupervised, closed-form linear transformation that aligns the second-order statistics (covariance matrices) of the source and target feature spaces. **Crucially, CORAL requires ZERO labels from the target network**, making it completely deployable in live production networks.

■ MATHEMATICAL FORMULATION OF CORAL ALIGNMENT:

1. Calculate source covariance matrix: $\mathbf{C}_s = \text{Cov}(\mathbf{X}_{\text{source}}) + \lambda \cdot \mathbf{I}$
 2. Calculate target covariance matrix: $\mathbf{C}_t = \text{Cov}(\mathbf{X}_{\text{target}}) + \lambda \cdot \mathbf{I}$
 3. Whiten source data and recolor to target: $\mathbf{X}_{\text{aligned}} = \mathbf{X}_{\text{source}} \cdot (\mathbf{C}_s)^{-1/2} \cdot (\mathbf{C}_t)^{1/2}$
- Result: The source training data now shares the exact geometric covariance orientation of the target testbed.*

Part 4: Quantifying the Recovery — Generalization Recovery Ratio (GRR)

To scientifically validate how effectively our engineering remediation closes the generalization gap, we formulate the **Generalization Recovery Ratio (GRR)**:

$$\text{GRR} = [\text{Score}(\text{aligned cross}) - \text{Score}(\text{raw cross})] / [\text{Score}(\text{same dataset}) - \text{Score}(\text{raw cross})]$$

- **Score(same dataset):** The laboratory baseline ceiling (e.g., 99% accuracy within CICIDS).
- **Score(raw cross):** The collapsed zero-shot transfer score before alignment (e.g., 45% on UNSW).
- **Score(aligned cross):** The recovered score after our feature selection & CORAL alignment (e.g., 78%).

Example Interpretation: $\text{GRR} = (78 - 45) / (99 - 45) = 33 / 54 = 61.1\%$ of the generalization deficit recovered.

Part 5: The Final Deliverable — Streamlit SOC Security Dashboard

To make our research actionable for real-world Security Operations Center (SOC) engineers, all models, divergence metrics, and alignment routines are packaged into an interactive web application built with **Streamlit and Plotly**:

- **Interactive Model Transfer Sandbox:** Select any source dataset (CICIDS or UNSW), select any target dataset, pick from 5 ML algorithms (DT, RF, LR, SVM, XGBoost), and toggle CORAL alignment On/Off.
- **Dynamic Confusion Matrix Heatmaps:** Watch False Alarm Rates drop in real time when domain alignment is activated.
- **Comparative ROC & Precision-Recall Curves:** Visual overlay comparing baseline ROC vs unaligned transfer ROC vs aligned transfer ROC.
- **Feature Distribution Shift Inspector:** Interactive Plotly overlays displaying the feature distribution shift before and after CORAL transformation.
- **Automated Transferability Report:** One-click export summarizing model degradation, GRR scores, and recommended deployment guidelines.

Part 6: Project Novelty vs. Prior Literature (Examiner Defense Matrix)

Aspect	Verkerken et al. (2022) / Cantone et al. (2024)	Our Project (USTM 2026)
Primary Goal	Observational diagnosis (documenting that models fail across datasets)	Active engineering remediation (diagnosing <i>and</i> fixing the collapse)
Feature Alignment	Stated 'harmonization is mandatory', but did not build an ontology	Built full semantic ontology reconciling CICFlowMeter ↔ Argus/Zeek
Domain Adaptation	None attempted (evaluated raw unaligned features only)	Implemented unsupervised CORAL covariance alignment
Mathematical Diagnosis	Only reported final accuracy drops (28%–42%)	Calculates Wasserstein Distance, KL Divergence & MMD per feature
Recovery Metric	None (no recovery attempted)	Formulated Generalization Recovery Ratio (GRR) metric
Delivery Format	Static research paper tables and static charts	Production-ready, interactive Streamlit SOC security dashboard